FREE · OPEN · PRACTICAL

AI Safety Playbook for Small Businesses

A free, practical guide to using AI tools safely, with templates you can copy.

By Emmanuel Diserbeau · Version 1.0 (draft) · October 2026

> **Not legal advice.** This playbook is general guidance for small businesses. It is not legal, security, or compliance advice. Laws differ by state, country, and industry. If you handle health, financial, children's, or government data, or you work in a regulated industry, talk to a qualified lawyer or advisor before you finalize your policy. The policy template in this playbook is a generic starting point. It is not the policy of any specific company.

Who this is for: owners and managers of businesses with 5 to 200 people who don't have a legal or security team.

How to use it: Read Section 1 (5 minutes). Fill in the policy in Section 2 (1 to 2 hours). Score yourself with Section 3. Pick a few safe wins from Section 4. Follow the 30-day plan in Section 5. Print Section 6 and hand it to every employee.

> **License:** Licensed under CC BY 4.0 (https://creativecommons.org/licenses/by/4.0/). You may share and adapt this material for any purpose, including commercially, with attribution to Emmanuel Diserbeau. **Suggested attribution:** "AI Safety Playbook for Small Businesses" by Emmanuel Diserbeau, licensed under CC BY 4.0 (https://creativecommons.org/licenses/by/4.0/). If you change it, add: "Adapted by [YOUR NAME OR COMPANY]."

Contents

1	Why AI Safety Matters	2
2	AI Acceptable Use Policy (Template)	3
3	AI Risk Self-Assessment	6
4	High-ROI, Low-Risk Use Cases by Department	7
5	30-Day Safe Rollout Plan	10
6	Employee Quick-Reference Card	11
7	References and Further Reading	12

1. Why AI Safety Matters

AI tools like ChatGPT, Grok, Gemini, Claude, and Copilot can save your team hours every week. They draft emails, summarize documents, write formulas, and answer questions in seconds. Your people are probably already using them, with or without permission.

That's the point of this playbook. You don't need to ban AI. You need a few simple rules so that AI helps the business instead of hurting it. Here are the five risks that matter most.

1. Data leaks

Anything you type or upload into an AI tool leaves your building. On free or personal accounts, some vendors may keep it, let staff review it, or use it to train future models. One pasted customer list or contract can become a privacy problem, a breach of contract, or a regulatory issue.

2. Intellectual property (IP)

Your pricing, processes, designs, source code, and customer lists are what make you competitive. Pasting them into the wrong tool can weaken trade-secret protection. On the other side, AI output can resemble someone else's copyrighted work. And in the U.S., content made purely by AI generally can't be copyrighted, so you may not "own" it the way you think.

3. Accuracy and "hallucinations"

AI tools sound confident even when they're wrong. They can invent facts, numbers, quotes, legal citations, and product specs. If nobody checks, those mistakes end up in quotes, contracts, and customer emails with your name on them.

4. Bias and fairness

AI learns from human data, so it can repeat human bias. That matters most when AI touches decisions about people: hiring, firing, pay, credit, or pricing. Using AI to screen résumés or rank employees can create discrimination risk, even if nobody intended it.

5. Customer trust

Customers forgive honest mistakes. They're less forgiving when they feel tricked. A fake-sounding "personal" message, an AI chatbot that gives wrong answers, or a leak of their data can undo years of goodwill.

> **The good news:** Most of these risks are handled by five habits. Use approved business accounts. Never paste sensitive data. Have a human check the output. Be honest with customers. Report problems fast. The rest of this playbook turns those habits into a policy, a checklist, and a plan.

2. AI Acceptable Use Policy (Template)

Copy this template into your own document. Replace everything in [BRACKETS]. Delete anything that doesn't fit your business. Keep it short enough that people will actually read it.

[COMPANY NAME] AI Acceptable Use Policy

Effective date: [DATE] · **Policy owner:** [NAME, TITLE] · **Next review:** [DATE]

2.1 Purpose and scope

[COMPANY NAME] encourages employees to use AI tools to work smarter. This policy explains how to do that without putting our customers, our people, or our business at risk. It applies to all employees, contractors, and temps who use AI tools for [COMPANY NAME] work, on any device, including personal phones.

"AI tools" means any software that generates text, images, audio, video, or code, or that analyzes data using AI. Examples include chatbots, writing assistants, meeting note-takers, and AI features built into software we already use.

2.2 Approved tools and accounts

1. Use only the AI tools on our **Approved AI Tools List**: [LIST TOOLS, e.g. "ChatGPT Business," "Microsoft 365 Copilot," "Grok Business"]. The current list lives at [LOCATION].
2. Use your **company-provided business account**, signed in with your work email. Do not use personal or free accounts for company work.
3. Where the vendor allows, our admin will **turn off training on our data** and set the shortest practical data-retention period. Do not change these settings.
4. Use multi-factor authentication (MFA) on every AI account.
5. Do not install AI browser extensions, plug-ins, "agents," or meeting note-takers unless they are on the approved list.
6. Do not connect AI tools to company email, files, or systems (for example, by granting access to a drive or inbox) without approval from [AI OWNER].

2.3 Data classification: what you can share

Class	What it is	Examples	OK in approved AI tools?
Public	Already public, or meant to be	Website text, published brochures, press releases	**Yes**
Internal	Everyday business info, low harm if leaked	General procedures, meeting agendas, non-sensitive drafts	**Yes**, in approved business accounts only
Confidential	Would hurt us or others if leaked	Pricing, margins, unreleased plans, vendor terms, customer names	**Only** if [AI OWNER] has approved that tool for Confidential data. Remove names and numbers when you can.
Restricted	Serious legal, financial, or safety harm if leaked	Everything on the "Never paste" list below	**Never**

When in doubt, treat it as one level higher, and ask [AI OWNER].

2.4 The "Never paste" list

Never type, paste, upload, or dictate any of the following into an AI tool, unless [AI OWNER] has given written approval for a specific tool and purpose:

- Customer personal information (names with contact details, addresses, ID numbers, account numbers)
- Financial records (bank details, card numbers, non-public financial statements, payroll)
- Passwords, API keys, access tokens, security codes, or system configurations
- Contracts, NDAs, legal documents, and anything a lawyer sent us
- Source code, scripts, and technical designs that belong to us or a client
- Trade secrets (formulas, processes, pricing models, supplier lists, strategy)
- Health or medical information about anyone
- Employee records (performance reviews, discipline, pay, benefits, medical leave, background checks)
- Anything we've promised a customer or partner to keep confidential

Tip: You can often get the same help by replacing real details with placeholders, such as "Customer A," "$X," or "Product 1."

2.5 Human review of AI output

1. You are responsible for anything you create with AI. "The AI wrote it" is never an excuse.
2. Check facts, numbers, names, dates, links, and citations before you use or send AI output.
3. A qualified person must review AI output before it is used for: anything sent to customers, legal or contract language, financial figures, safety or technical specifications, and code that will run on our systems.
4. AI must not make final decisions about people (hiring, firing, pay, promotion, discipline, credit). A person decides, and can explain why.

2.6 Being honest with customers

1. Don't present AI-generated content as a personal, handwritten, or expert-reviewed message when it isn't.
2. If customers talk to an AI chatbot or voice agent, tell them it's AI and give them an easy way to reach a person.
3. Never use AI to create fake reviews, testimonials, or images of real people.
4. Follow any customer contract that limits or requires disclosure of AI use.

2.7 Intellectual property and copyright

1. Don't ask AI to copy or imitate a specific copyrighted work, brand, logo, or living artist's style for commercial use.
2. Don't upload other people's copyrighted material (books, paid reports, competitor content) unless we have the right to use it that way.
3. Assume purely AI-generated content may not be protected by copyright. For important work (logos, product names, marketing campaigns), add meaningful human creative input, and check with [AI OWNER] before relying on it.
4. Check that AI-written code doesn't include licensed code we can't use.

2.8 Approving new AI tools

Before anyone uses a new AI tool or feature for company work, send a request to [AI OWNER] at [EMAIL/FORM]. We will check:

- Does the vendor use our data for training, and can we turn it off?

- How long is data kept, and where?
- Does it offer business accounts, admin controls, and MFA?
- Does it meet any customer contract or industry requirements we have?
- What's the real business benefit?

We aim to answer within [5] business days.

2.9 Reporting problems

Report it to [AI OWNER / CONTACT] within [24 hours] if you:

- Pasted something you shouldn't have
- See AI output that is offensive, biased, or seriously wrong and it reached a customer
- Notice an AI tool behaving strangely or asking for unusual access
- Spot someone using unapproved tools with company data

We won't punish honest, prompt reports of mistakes. Fast reporting lets us fix problems while they're small.

2.10 Training

Everyone must complete [COMPANY NAME]'s AI basics training (about [30] minutes) before using AI tools for work, and a short refresher every [12] months.

2.11 Consequences

Breaking this policy can lead to loss of AI tool access and disciplinary action, up to and including termination, in line with our other policies. Intentional misuse of Restricted data will be treated seriously.

2.12 Review

[AI OWNER] will review this policy every [6] months, or sooner if laws, tools, or our business change.

2.13 Employee acknowledgment

I have read and understand the [COMPANY NAME] AI Acceptable Use Policy. I agree to follow it and to ask when I'm not sure.

Employee name	Signature	Date

3. AI Risk Self-Assessment

Answer each question honestly. Score **0** for Yes, **1** for Partly or Not sure, and **2** for No. Add up your score.

#	Question	Yes (0)	Partly (1)	No (2)
1	Do you have a written AI use policy that every employee has read and signed?	☐	☐	☐
2	Do you know which AI tools your staff actually use today, including on phones?	☐	☐	☐
3	Do staff use business accounts (not personal or free ones) for company work?	☐	☐	☐
4	Have you turned off training on your data and set data retention, where vendors allow?	☐	☐	☐
5	Do employees know what data must never be pasted into AI tools?	☐	☐	☐
6	Is AI output checked by a person before it reaches customers or affects money, contracts, or safety?	☐	☐	☐
7	Is AI kept out of final decisions about hiring, pay, discipline, or credit?	☐	☐	☐
8	Is there one named person who owns AI decisions and approves new tools?	☐	☐	☐
9	Do employees know how to report an AI mistake or data slip, without fear of blame?	☐	☐	☐
10	Are MFA and an offboarding step (remove access when people leave) in place for AI accounts?	☐	☐	☐

Your total score: _____ / 20

Score	Risk level	What it means	What to do
0–5	**LOW**	You have the basics covered.	Keep it up. Review the policy every 6 months. Expand to more use cases from Section 4. Re-score each quarter.
6–12	**MEDIUM**	Some gaps that could cause a leak or a costly mistake.	Fix every question you scored 2 within 30 days. Start with questions 3, 5, and 6. Follow the 30-day plan in Section 5.
13–20	**HIGH**	Staff are likely using AI with no guardrails.	Act this week. Name an AI owner. Send a short interim "never paste" notice to all staff today. Buy business accounts for the tools people already use. Then follow Section 5.

4. High-ROI, Low-Risk Use Cases by Department

Start where the payoff is high and the risk is low. Each idea is rated:

- **SAFE:** Good to go with approved tools and normal human review.
- **CAUTION:** Fine with care. Remove sensitive details, and have a qualified person check the result.
- **AVOID:** Don't use AI this way. A safer alternative is shown.

Prompts use [brackets] for details you fill in. Never put "Never paste" data in the brackets.

Sales

Use case	Rating	Example prompt
Draft a follow-up email after a meeting	SAFE	"Write a friendly, 120-word follow-up email to a [industry] prospect after a first meeting. Thank them, recap that we discussed [topic], and propose a call next week."
Prep for a sales call from public info	SAFE	"Using only public information, list 5 smart questions to ask a [industry] company with about [size] employees that is looking at [product type]."
Draft a proposal from your notes	CAUTION	"Turn these notes into a one-page proposal outline with sections for problem, solution, timeline, and next steps: [notes with customer names and prices removed]."
Upload your CRM export to "find the best leads"	AVOID	Instead: "What criteria should a small [industry] business use to score sales leads? Give me a simple 1–5 scoring sheet."

Operations

Use case	Rating	Example prompt
Write a standard operating procedure (SOP)	SAFE	"Turn these rough steps into a clear SOP with a purpose, numbered steps, and a safety checklist: [steps]."
Build a training checklist for new hires	SAFE	"Create a first-week checklist for a new [role] at a [type of] business, grouped by day."
Analyze vendor or inventory trends	CAUTION	"Here is a table of monthly order volumes by product category, with vendor names removed: [data]. Point out trends and anything unusual."
Let AI write safety or technical specs with no expert review	AVOID	Instead: "List the questions I should ask a qualified engineer when reviewing a [type of] specification."

Finance

Use case	Rating	Example prompt
Explain or write a spreadsheet formula	SAFE	"Write an Excel formula that sums column D when column B says 'Open' and column C is older than 30 days. Explain how it works."
Draft a plain-English summary of a finance concept for staff	SAFE	"Explain working capital to non-finance managers in 150 words, with one simple example."
Draft a polite payment-reminder template	CAUTION	"Write a firm but friendly 3-step payment reminder sequence for invoices 15, 30, and 45 days overdue. Use [CUSTOMER] and [AMOUNT] as placeholders."

Use case	Rating	Example prompt
Paste bank statements, payroll, or tax records	AVOID	Instead: "What checks should a small business do each month when reconciling bank accounts? Make it a checklist."

Human Resources (HR)

Use case	Rating	Example prompt
Draft a job description	SAFE	"Write an inclusive job description for a [role] at a [size] [industry] company. Use plain language and avoid biased wording. Include must-haves vs nice-to-haves."
Create interview questions	SAFE	"Create 8 structured, job-related interview questions for a [role], with what a strong answer looks like."
Draft a policy or handbook section	CAUTION	"Draft a plain-English remote work policy for a 40-person company. Flag anything that may depend on [state] law." (Have HR or a lawyer review it.)
Screen or rank candidates or employees with AI	AVOID	Instead, use AI to write a fair scoring rubric, and have people do the scoring.

Customer Service

Use case	Rating	Example prompt
Build reply templates for common questions	SAFE	"Write 5 short, friendly reply templates for: order status, return request, shipping delay, price question, and complaint. Leave [placeholders] for details."
Turn your FAQ into clearer help articles	SAFE	"Rewrite this public FAQ answer so a customer can understand it in 30 seconds: [public FAQ text]."
Draft a reply to an upset customer	CAUTION	"Help me reply calmly to a customer who is upset about [issue, no names or account details]. Apologize, explain next steps, and offer [remedy]."
Let a chatbot answer customers with no disclosure or human backup	AVOID	Instead: label any bot as AI, limit it to approved answers, and give customers an easy path to a person.

IT

Use case	Rating	Example prompt
Write user how-to guides	SAFE	"Write a step-by-step guide for employees to set up multi-factor authentication on [app], for non-technical readers."
Draft a security awareness tip	SAFE	"Write a 100-word monthly security tip about spotting phishing emails that use AI-written text."
Get help with scripts or troubleshooting	CAUTION	"Explain this error message and suggest fixes: [error text with server names, IPs, and usernames removed]." (Test in a safe environment first.)
Paste passwords, keys, configs, or proprietary code	AVOID	Instead: "What are best practices for storing API keys in a small company?" Use a password manager or secrets vault.

5. 30-Day Safe Rollout Plan

This plan assumes one owner who can spend 2 to 4 hours a week. Adjust the pace to fit your business.

Week 1: Get the basics in place

- ☐ Name an **AI owner** (often the IT lead, operations manager, or owner).
- ☐ Take the **Self-Assessment** in Section 3 and note your score.
- ☐ Ask staff (anonymously, if it helps) which AI tools they already use and for what.
- ☐ Send an **interim notice** to all staff: "Until our policy is ready, don't paste customer, financial, employee, or confidential data into any AI tool."
- ☐ Pick 1 or 2 tools to approve. Choose business plans with admin controls.

Week 2: Set up and write the rules

- ☐ Buy business accounts. Turn on MFA. Turn off training on your data where allowed. Set data retention.
- ☐ Fill in the **AI Acceptable Use Policy** template in Section 2.
- ☐ Publish the **Approved AI Tools List** and the request process for new tools.
- ☐ Pick 3 to 5 **pilot users** from different departments.
- ☐ Choose 3 to 5 **SAFE** use cases from Section 4 for the pilot.

Week 3: Train and pilot

- ☐ Run a 30-minute **training session**: the policy, the "Never paste" list, human review, and how to report problems.
- ☐ Hand out the **Quick-Reference Card** (Section 6).
- ☐ Pilot users try the chosen use cases and track time saved, quality, and any issues.
- ☐ Hold a 15-minute check-in at the end of the week. Fix anything confusing.

Week 4: Roll out and measure

- ☐ Collect signed **acknowledgments** from everyone who will use AI.
- ☐ Open approved tools to all trained staff.
- ☐ Share 2 or 3 pilot wins (with real time savings) to build momentum.
- ☐ Re-take the **Self-Assessment** and compare scores.
- ☐ Put the next policy review on the calendar (every 6 months), and plan which CAUTION use cases to try next.

> **Measure what matters:** Track hours saved per week, the number of people using approved tools, the number of incidents reported, and one quality measure (for example, customer reply time). Those four numbers tell you if AI is paying off safely.

6. Employee Quick-Reference Card

Print this page and keep it at your desk.

AI at [COMPANY NAME]: Do's and Don'ts

✓ DO	✗ DON'T
Use only **approved tools** with your **work account**	Use personal or free AI accounts for company work
Replace real details with **placeholders** ("Customer A," "$X")	Paste customer info, financials, passwords, contracts, code, trade secrets, health data, or employee records
Check every fact, number, name, and link before using AI output	Send AI output to a customer without reading it
Get a **qualified review** for legal, financial, safety, or technical content	Let AI make final decisions about people (hiring, pay, discipline)
Be **honest** with customers when AI is involved	Create fake reviews, testimonials, or images of real people
Ask [AI OWNER] before trying a **new tool**, plug-in, or note-taker	Connect AI to email, files, or systems without approval
Report mistakes within 24 hours. No blame for honest reports.	Hide a mistake or hope nobody notices

Before you hit Enter, ask yourself:

1. Is this an **approved tool** and my **work account**?
2. Would I be OK if this text were **posted on a billboard** with our name on it?
3. Will a **person check** the answer before it's used?

Questions or problems? Contact [AI OWNER NAME] at [EMAIL / PHONE / CHAT CHANNEL].

7. References and Further Reading

These free, well-known resources go deeper. All links were checked in October 2026.

Risk management frameworks

- **NIST AI Risk Management Framework (AI RMF 1.0).** U.S. National Institute of Standards and Technology, January 2023. A voluntary framework built around four functions: Govern, Map, Measure, and Manage. https://www.nist.gov/itl/ai-risk-management-framework
- **NIST AI 600-1: Generative AI Profile.** NIST, July 2024. A companion to the AI RMF focused on risks of generative AI tools. https://doi.org/10.6028/NIST.AI.600-1
- **NIST AI RMF Playbook and AI Resource Center.** Suggested actions for putting the framework into practice. https://airc.nist.gov/
- **ISO/IEC 42001:2023, AI management systems.** The international standard for managing AI in an organization (paid standard). https://www.iso.org/standard/42001
- **OECD AI Principles.** The first intergovernmental standard on trustworthy AI, adopted 2019 and updated 2024. https://www.oecd.org/en/topics/sub-issues/ai-principles.html

Security guidance

- **OWASP Top 10 for Large Language Model Applications.** Community list of top AI security risks, such as prompt injection and sensitive information disclosure. https://genai.owasp.org/llm-top-10/
- **Deploying AI Systems Securely.** Joint guidance from NSA, CISA, FBI, and international partners, April 2024. https://www.cisa.gov/news-events/alerts/2024/04/15/joint-guidance-deploying-ai-systems-securely
- **AI Data Security: Best Practices for Securing Data Used to Train and Operate AI Systems.** CISA and partners, May 2025. https://www.cisa.gov/resources-tools/resources/ai-data-security-best-practices-securing-data-used-train-operate-ai-systems
- **Guidelines for Secure AI System Development.** UK NCSC, CISA, and partners, 2023. https://www.ncsc.gov.uk/files/Guidelines-for-secure-AI-system-development.pdf

Consumer protection, copyright, and law

- **U.S. Federal Trade Commission (FTC): Artificial Intelligence.** News, enforcement actions, and policy statements on AI, including cases about deceptive AI claims and fake reviews. https://www.ftc.gov/industry/technology/artificial-intelligence
- **U.S. Copyright Office: Copyright and Artificial Intelligence.** Reports on AI and copyright, including Part 2 on copyrightability (January 2025). https://www.copyright.gov/ai/
- **EU Artificial Intelligence Act (Regulation (EU) 2024/1689).** Relevant if you sell to or operate in the EU. https://op.europa.eu/en/publication-detail/-/publication/dc8116a1-3fe6-11ef-865a-01aa75ed71a1/language-en

AI Safety Playbook for Small Businesses, v1.0 (draft), by Emmanuel Diserbeau. Licensed under CC BY 4.0 (https://creativecommons.org/licenses/by/4.0/). You may share and adapt this material for any purpose, including commercially, with attribution to Emmanuel Diserbeau. Not legal advice.